

University of Central Florida

CGS 2545

Database Concepts

DEPARTMENT OF ELECTRICAL ENGINEERING & COMPUTER SCIENCE
COMPUTER SCIENCE DIVISION

Concurrency Control

- In a multiprogramming environment where multiple transactions can be executed simultaneously, it is highly important to control the concurrency of transactions.
- We have concurrency control protocols to ensure atomicity, isolation, and serializability of concurrent transactions.
- Concurrency control protocols can be broadly divided into two categories
 - Lock based protocols
 - Time stamp based protocols

Concurrency Control

- Lock-based Protocols
 - Database systems equipped with lock-based protocols use a mechanism by which any transaction cannot read or write data until it acquires an appropriate lock on it.
 - Locks are of two kinds
 - Binary Locks
 - Shared/exclusive

Concurrency Control

- Lock-based Protocols
 - **Binary Locks**
 - A lock on a data item can be in two states; it is either locked or unlocked.
 - **Shared/exclusive**
 - This type of locking mechanism differentiates the locks based on their uses.
 - If a lock is acquired on a data item to perform a write operation, it is an exclusive lock.
 - Allowing more than one transaction to write on the same data item would lead the database into an inconsistent state.
 - Read locks are shared because no data value is being changed.

Concurrency Control

- Lock-based Protocols
 - There are four types of lock protocols available
 - Simplistic Lock Protocol
 - Pre-claiming Lock Protocol
 - Two-Phase Locking 2PL
 - Strict Two-Phase Locking

Concurrency Control

- Lock-based Protocols
 - Simplistic Lock Protocol
 - Simplistic lock-based protocols allow transactions to obtain a lock on every object before a 'write' operation is performed.
 - Transactions may unlock the data item after completing the 'write' operation.

Concurrency Control

- Lock-based Protocols
 - Pre-claiming Lock Protocol
 - Pre-claiming protocols evaluate their operations and create a list of data items on which they need locks.
 - Before initiating an execution, the transaction requests the system for all the locks it needs beforehand.
 - If all the locks are granted, the transaction executes and releases all the locks when all its operations are over.
 - If all the locks are not granted, the transaction rolls back and waits until all the locks are granted.


Concurrency Control

- Lock-based Protocols
 - Pre-claiming Lock Protocol
 - Pre-claiming protocols evaluate their operations and create a list of data items on which they need locks.
 - Before initiating an execution, the transaction requests the system for all the locks it needs beforehand.
 - If all the locks are granted, the transaction executes and releases all the locks when all its operations are over.
 - If all the locks are not granted, the transaction rolls back and waits until all the locks are granted.

Concurrency Control

- Lock-based Protocols
 - Pre-claiming Lock Protocol

Lock acquisition
phase




Concurrency Control

- Lock-based Protocols
 - Two-Phase Locking 2PL
 - This locking protocol divides the execution phase of a transaction into three parts.
 - In the first part, when the transaction starts executing, it seeks permission for the locks it requires.
 - The second part is where the transaction acquires all the locks.
 - As soon as the transaction releases its first lock, the third phase starts. In this phase, the transaction cannot demand any new locks; it only releases the acquired locks.

Concurrency Control

- Lock-based Protocols
 - Two-Phase Locking 2PL



Concurrency Control


- Lock-based Protocols
 - Two-Phase Locking 2PL
 - Two-phase locking has two phases, one is **growing**, where all the locks are being acquired by the transaction; and the second phase is shrinking, where the locks held by the transaction are being released.
 - To claim an exclusive (write) lock, a transaction must first acquire a shared (read) lock and then upgrade it to an exclusive lock.

Concurrency Control

- Lock-based Protocols
 - Strict Two-Phase Locking
 - The first phase of Strict-2PL is same as 2PL.
 - After acquiring all the locks in the first phase, the transaction continues to execute normally.
 - But in contrast to 2PL, Strict-2PL does not release a lock after using it.
 - Strict-2PL holds all the locks until the commit point and releases all the locks at a time.
 - Strict-2PL does not have cascading abort as 2PL does

Concurrency Control

- Lock-based Protocols
 - Strict Two-Phase Locking



Concurrency Control

- Timestamp-based Protocols
 - The most commonly used concurrency protocol is the timestamp based protocol.
 - This protocol uses either system time or logical counter as a timestamp.
 - Lock-based protocols manage the order between the conflicting pairs among transactions at the time of execution, whereas timestamp-based protocols start working as soon as a transaction is created.

Concurrency Control

- Timestamp-based Protocols
 - Every transaction has a timestamp associated with it, and the ordering is determined by the age of the transaction.
 - A transaction created at 0002 clock time would be older than all other transactions that come after it.
 - For example, any transaction 'y' entering the system at 0004 is two seconds younger and the priority would be given to the older one.

Concurrency Control

- Timestamp-based Protocols
 - In addition, every data item is given the latest read and write-timestamp.
 - This lets the system know when the last 'read and write' operation was performed on the data item.

Concurrency Control

- Timestamp Ordering Protocol
 - The timestamp-ordering protocol ensures serializability among transactions in their conflicting read and write operations.
 - This is the responsibility of the protocol system that the conflicting pair of tasks should be executed according to the timestamp values of the transactions.
 - The timestamp of transaction T_i is denoted as $TS(T_i)$
 - Read time-stamp of data-item X is denoted by $R\text{-timestamp}(X)$.
 - Write time-stamp of data-item X is denoted by $W\text{-timestamp}(X)$.

Concurrency Control

- Timestamp Ordering Protocol

- Timestamp ordering protocol works as follows

If a transaction T_i issues a read(X) operation –

- ▣ If $TS(T_i) < W\text{-timestamp}(X)$
 - ▣ Operation rejected.
- ▣ If $TS(T_i) \geq W\text{-timestamp}(X)$
 - ▣ Operation executed.
- ▣ All data-item timestamps updated.

If a transaction T_i issues a write(X) operation –

- ▣ If $TS(T_i) < R\text{-timestamp}(X)$
 - ▣ Operation rejected.
- ▣ If $TS(T_i) < W\text{-timestamp}(X)$
 - ▣ Operation rejected and T_i rolled back.
- ▣ Otherwise, operation executed.

Concurrency Control

- Timestamp Ordering Protocol

Thomas' Write Rule

This rule states if $TS(T_i) < W\text{-timestamp}(X)$, then the operation is rejected and T_i is rolled back.

Time-stamp ordering rules can be modified to make the schedule view serializable.

Instead of making T_i rolled back, the 'write' operation itself is ignored.